

Hyperparameter Optimization

using  **DASK**

with  **Orion**



DASK
Distributed Summit 2021

Xavier
Bouthillier



Yonggang
Hu




Agenda

- Introduction to black-box optimization


- Introduction to  Orion

- Integration with  DASK




Optimization

Optimization




Learning well?

Optimization




dask.org




Faster?

Optimization






Optimization Search Spaces






Optimization Methods

Grid Search



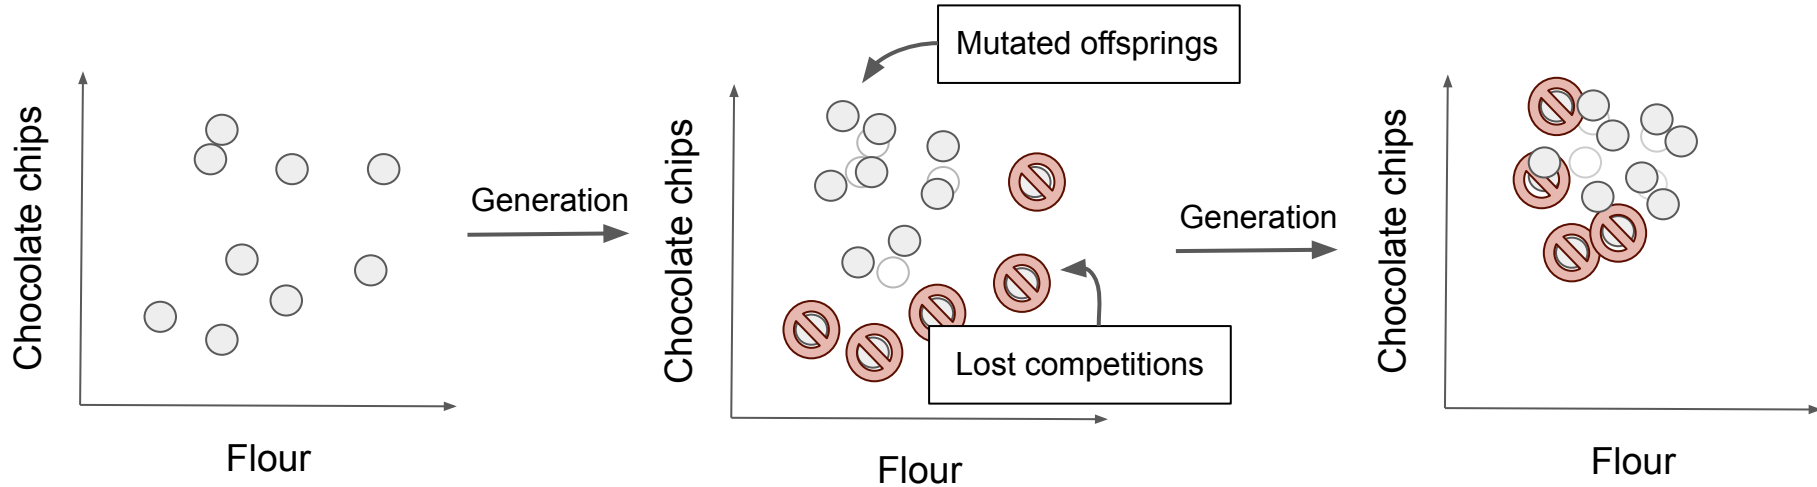
Random Search





Optimization Methods

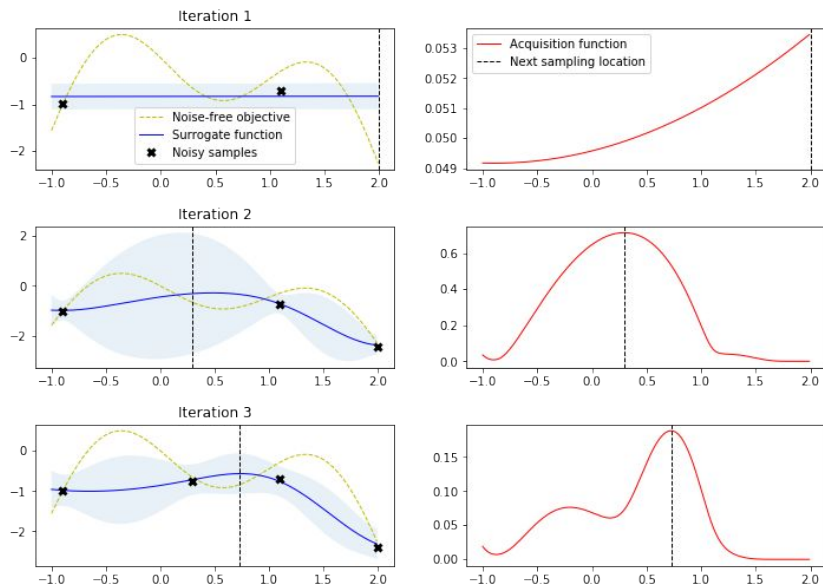
Evolutionary Algorithms





Optimization Methods

Model-Based Algorithms




krasserm.github.io/2018/03/21/bayesian-optimization/

Design Principles of

The trigger: Finicky HPC

- No nodes for persistent processes.
- No internet connection on compute nodes.
- Unstable, you would often end up on nodes with faulty GPUs.
- Very slow file-system. Maybe due to bad usage, but whatever the reason, it was terribly slow. ͡_͡(ಠ)_͡

Design Principles of Orion

- Resiliency: Master-less and resilient to worker failures.
- Interoperability: Agnostic to programming language of the .
- Modularity: Most components are 'pluginable'.

Installing

From PyPI


```
$ pip install orion
```

From Anaconda

```
$ conda install -c epistimio orion
```

Running Orion

```
def rosenbrock(x, y, a=1, b=100):  
    z = (a - x)**2 + b * (y - x**2)**2  
  
    return [  
        {"name": "objective",  
         "type": "objective",  
         "value": z  
        }  
    ]
```



Running

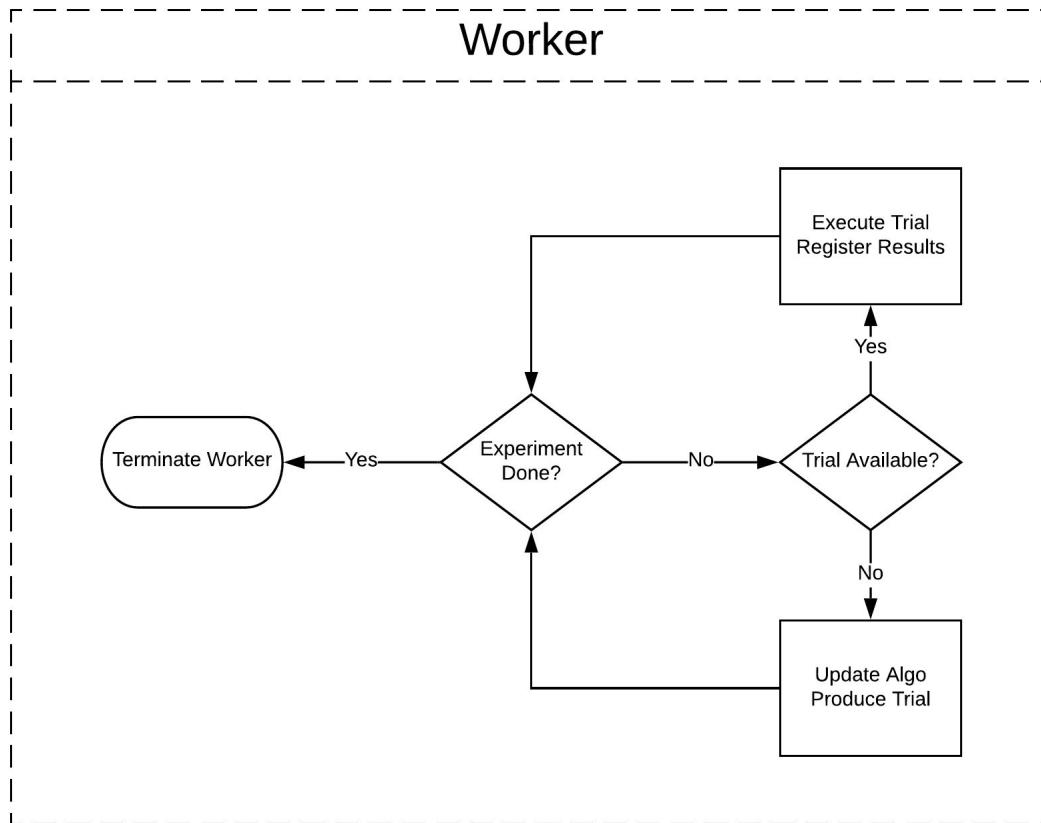
```
def rosenbrock(x, y, a=1, b=100):  
    z = (a - x)**2 + b * (y - x**2)**2  
  
    return [  
        {"name": "objective",  
         "type": "objective",  
         "value": z  
        }  
    ]
```

```
from orion.client import build_experiment
```

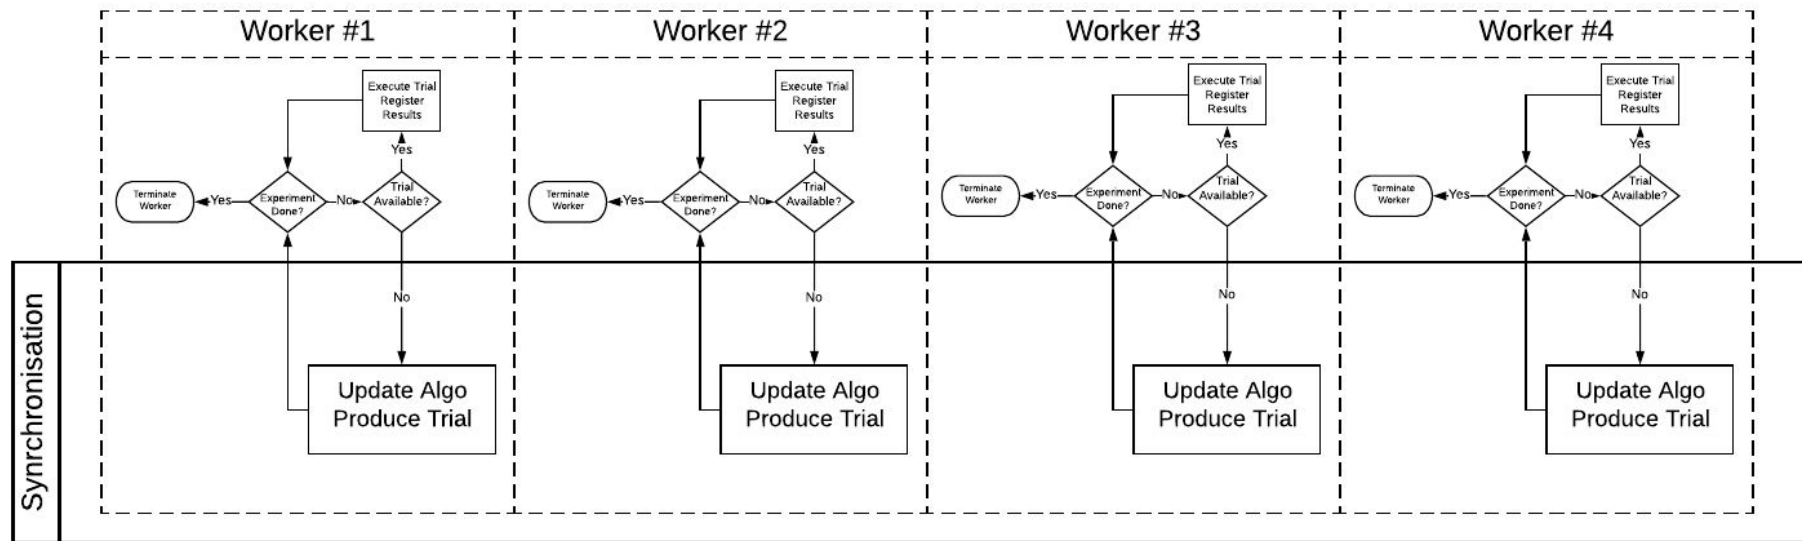
```
experiment = build_experiment(  
    "random-rosenbrock",  
    space={  
        "x": "uniform(-5, 5)",  
        "y": "uniform(-5, 5)",  
    },  
)
```

```
experiment.workon(rosenbrock, max_trials=20)
```


Workflow



Distributed Workflow



Running

```
def rosenbrock(x, y, a=1, b=100):  
    z = (a - x)**2 + b * (y - x**2)**2  
  
    return [  
        {"name": "objective",  
         "type": "objective",  
         "value": z  
        }  
    ]
```

```
from orion.client import build_experiment
```

```
experiment = build_experiment(  
    "random-rosenbrock",  
    space={  
        "x": "uniform(-5, 5)",  
        "y": "uniform(-5, 5)",  
    },  
)
```

```
experiment.workon(rosenbrock, max_trials=20)
```

Running Orion with DASK

```
def rosenbrock(x, y, a=1, b=100):  
    z = (a - x)**2 + b * (y - x**2)**2  
  
    return [  
        {"name": "objective",  
         "type": "objective",  
         "value": z  
        }  
    ]
```

```
from orion.client import build_experiment
```

```
experiment = build_experiment(  
    "random-rosenbrock",  
    space={  
        "x": "uniform(-5, 5)",  
        "y": "uniform(-5, 5)",  
    },  
)
```

```
with experiment.tmp_executor("dask"):  
    experiment.workon(rosenbrock, max_trials=20)
```

Running with DASK

```
def main(C, gamma, tol, class_weight):
    diabetes = datasets.load_diabetes()

    X = diabetes.data
    y = diabetes.target

    model = SVC(kernel="rbf", C=C, gamma=gamma,
                 tol=tol, class_weight=class_weight)

    with joblib.parallel_backend("dask"):
        cv_results = cross_validate(model, X, y, cv=5)

    accuracy = numpy.mean(cv_results["test_score"])
    error_rate = 1 - accuracy

    return [{"name": "test_error_rate",
               "type": "objective",
               "value": error_rate}]
```

Running Orion with DASK

```
def main(C, gamma, tol, class_weight):
    diabetes = datasets.load_diabetes()

    X = diabetes.data
    y = diabetes.target

    model = SVC(kernel="rbf", C=C, gamma=gamma,
                 tol=tol, class_weight=class_weight)

    with joblib.parallel_backend("dask"):
        cv_results = cross_validate(model, X, y, cv=5)

    accuracy = numpy.mean(cv_results["test_score"])
    error_rate = 1 - accuracy

    return [{"name": "test_error_rate",
               "type": "objective",
               "value": error_rate}]
```

```
def hpo(n_workers=16):
```


```
    experiment = create_experiment(
        name="dask",
        max_trials=200,
        space={
            "C": "loguniform(1e-6, 1e6, precision=None)",
            "gamma": "loguniform(1e-8, 1e8, precision=None)",
            "tol": "loguniform(1e-4, 1e-1, precision=None)",
            "class_weight": "choices([None, 'balanced'])",
        },
    )

    with experiment.tmp_executor("dask", n_workers=n_workers):
        experiment.workon(main, n_workers=n_workers // 2)

    experiment.plot.regret().show()
    experiment.plot.partial_dependencies(
        params=["C", "gamma", "tol"]).show()
```

```
if __name__ == "__main__":
    hpo()
```

Running Orion with DASK

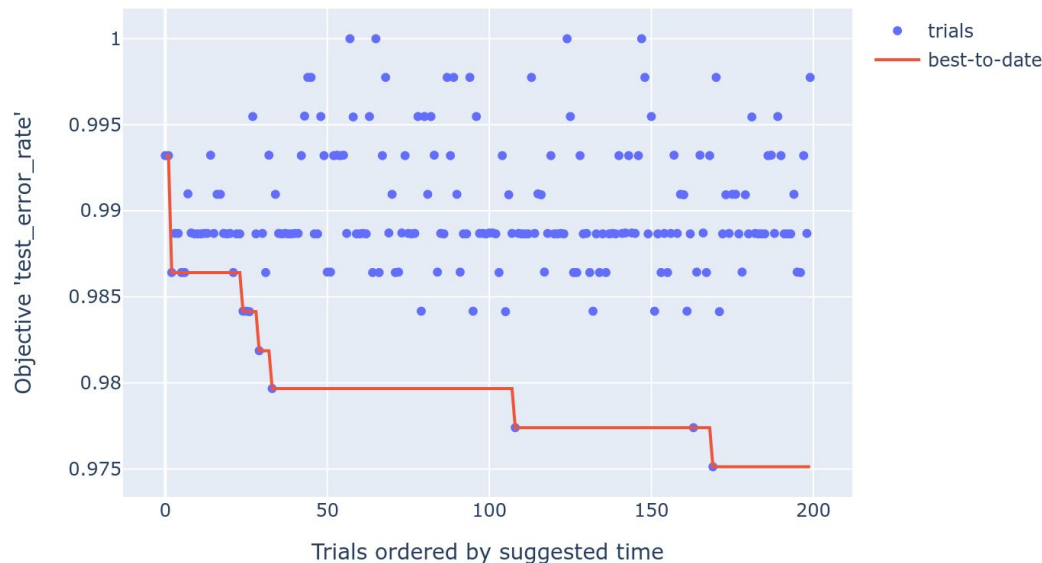


Running Orion with DASK

```
experiment.plot.regret().show()
```

Regret for experiment 'dask'

Conveys how quickly the algorithm found the best hyperparameters.



Running Orion with DASK

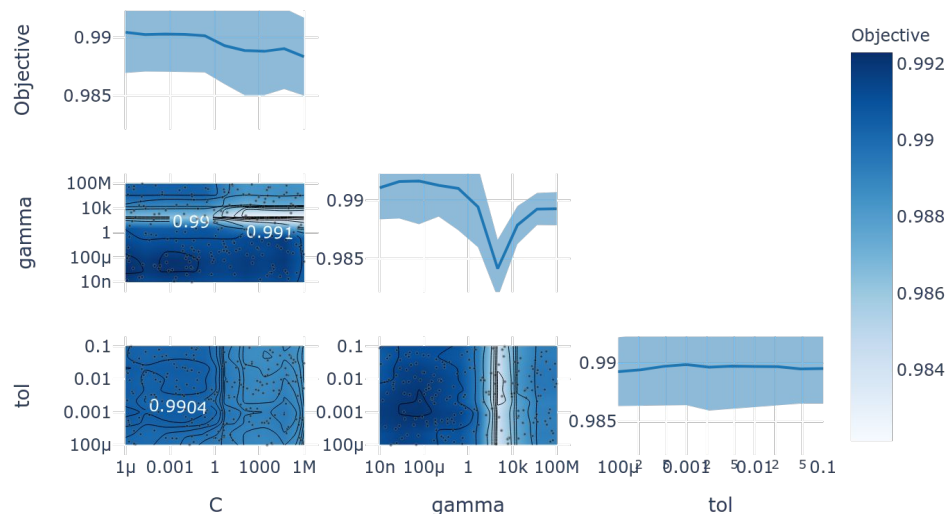
```
experiment.plot.partial_dependencies(  
    params=["C", "gamma", "tol"]).show()
```

Conveys overview of the search space, what has been explored.

Helps identifying best optimal regions of the space.



Partial dependencies for experiment 'dask'



Coming up next

Issues of master-less design

- Complexifies the implementation of algorithms. Lots of corner-cases!
- Requires more I/O with the database and scaling is hampered.
- Heavy reliance on atomicity of storage operations to handle race conditions.

Transition to Queue-based workflow

- Supports resiliency for the workers.
- Simpler to operate. Less asynchronicity to handle.
- More efficient. Less database I/O.
- Master becomes single point of failure.
- Needs persistent process on HPC.

Transition to Queue-based workflow

- Supports resiliency for the workers.
- Simpler to operate. Less asynchronicity to handle.
- More efficient. Less database I/O.
- Master becomes single point of failure.
- Needs persistent process on HPC.

**We need to have
the master
bouncing around
workers!**

Documentation

orion.readthedocs.io

Source code

github.com/Epistimio/orion

Packages

pypi.org/project/orion

anaconda.org/epistimio/orion

We are looking for interns and full time developers!

Send your CV to jobs@mila.quebec if you are interested.



References

Agnihotri, Apoorv, and Nipun Batra. "Exploring bayesian optimization." Distill 5, no. 5 (2020): e26.

Bergstra, James, Rémi Bardenet, Yoshua Bengio, and Balázs Kégl. "Algorithms for hyper-parameter optimization." In 25th annual conference on neural information processing systems (NIPS 2011), vol. 24. Neural Information Processing Systems Foundation, 2011.

Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." Journal of Machine Learning Research 13, no. Feb (2012): 281-305.

Biedenkapp, André, Joshua Marben, Marius Lindauer, and Frank Hutter. "Cave: Configuration assessment, visualization and evaluation." In International Conference on Learning and Intelligent Optimization, pp. 115-130. Springer, Cham, 2018.

Bouthillier, Xavier, and Gaël Varoquaux. "Survey of machine-learning experimental methods at NeurIPS2019 and ICLR2020." PhD diss., Inria Saclay Ile de France, 2020.

Eggersperger, Katharina, Marius Lindauer, and Frank Hutter. "Pitfalls and best practices in algorithm configuration." Journal of Artificial Intelligence Research 64 (2019): 861-893.

Henderson, P., R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger. "Deep Reinforcement Learning that Matters. arXiv, cs." (2017).

References

Hoffer, Elad, Ron Banner, Itay Golan, and Daniel Soudry. "Norm matters: efficient and accurate normalization schemes in deep networks." arXiv preprint arXiv:1803.01814 (2018).

Hutter, Frank, Lars Kotthoff, and Joaquin Vanschoren. Automated machine learning: methods, systems, challenges. Springer Nature, 2019.

Jastrzębski, Stanislaw, Zachary Kenton, Devansh Arpit, Nicolas Ballas, Asja Fischer, Yoshua Bengio, and Amos Storkey. "Width of minima reached by stochastic gradient descent is influenced by learning rate to batch size ratio." In International Conference on Artificial Neural Networks, pp. 392-402. Springer, Cham, 2018.

Van Laarhoven, Twan. "L2 regularization versus batch and weight normalization." arXiv preprint arXiv:1706.05350 (2017).

Li, Lisha, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. "Hyperband: A novel bandit-based approach to hyperparameter optimization." The Journal of Machine Learning Research 18, no. 1 (2017): 6765-6816.

Li, Liam, Kevin Jamieson, Afshin Rostamizadeh, Ekaterina Gonina, Moritz Hardt, Benjamin Recht, and Ameet Talwalkar. "Massively parallel hyperparameter tuning." arXiv preprint arXiv:1810.05934 (2018).

Li, Liam, and Ameet Talwalkar. "Random search and reproducibility for neural architecture search." In Uncertainty in Artificial Intelligence, pp. 367-377. PMLR, 2020.

References

Lindauer, Marius, and Frank Hutter. "Best practices for scientific research on neural architecture search." *Journal of Machine Learning Research* 21, no. 243 (2020): 1-18.

Snoek, Jasper, Hugo Larochelle, and Ryan P. Adams. "Practical bayesian optimization of machine learning algorithms." *arXiv preprint arXiv:1206.2944* (2012).

So, David, Quoc Le, and Chen Liang. "The evolved transformer." In *International Conference on Machine Learning*, pp. 5877-5886. PMLR, 2019.

Wilson, Ashia C., Rebecca Roelofs, Mitchell Stern, Nathan Srebro, and Benjamin Recht. "The marginal value of adaptive gradient methods in machine learning." *arXiv preprint arXiv:1705.08292* (2017).